

The recommended reading for this lesson is:

Zakir Durumeric, Eric Wustrow, and J. Alex Halderman. [ZMap: Fast Internet-Wide Scanning and its Security Applications](#).

In Proceedings of the 22nd USENIX Security Symposium

You can find short summary of the paper at the end of the document.

The Internet is a large ecosystem of networks. It is also a large ecosystem of vulnerabilities. To defend the network, we need to know it. In other words, we need to map the entire Internet. In this lesson, we will learn how to map the Internet efficiently and thoroughly, and survey vulnerabilities.



Attacker Intelligence Quiz

Attackers have three phases of intelligence gathering.
Match each phase to its description:

- B Footprinting (FP)
- A Scanning (S)
- C Enumeration (E)

- A. The attacker uses the Internet to obtain information on specific IP addresses. The kind of information gathered is: O.S., services, and architecture of the target system
- B. The attacker gathers information about a target. The kind of information gathered is: DNS, email servers, and the IP address range.
- C. The attacker gathers information on network user and group names, routing tables, and simple network management protocol.

information gathered is operating system, services, and architecture of the target system. In Enumeration, the attacker gathers information on network user and group names, routing tables and simple network management protocol.

Let us start with a quiz on how attackers can gather intelligence. Match each phrase with its description.

In Footprinting, the attacker gathers information about target. The kind of information gathered is DNS, email servers, and the IP address range. In Scanning, the attacker uses the Internet to obtain information on specific IP addresses. The kind of



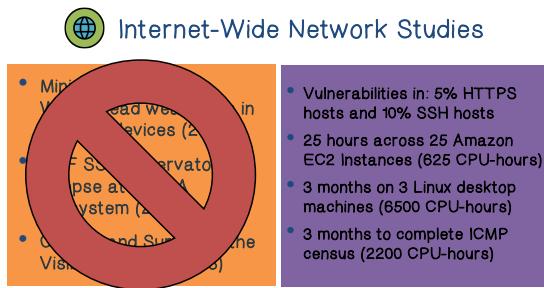
Internet Wide Security Scanning

- Expose new vulnerabilities
- Track adoption of defensive mechanisms
- Probing the ENTIRE address space with existing tools is both difficult and slow

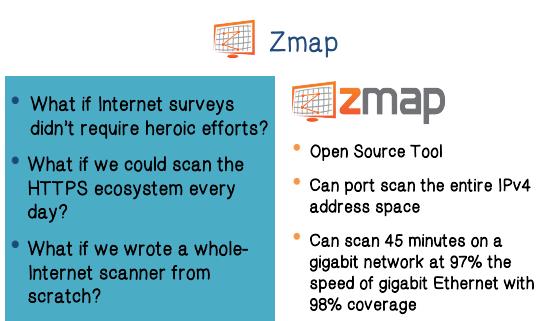
E.g., Datasets of cryptographic keys + the entire address space =
Detection of vulnerabilities and understanding of the ecosystem

There are a lot of benefits from Internet Wide Security Scanning. For example, we can discover new vulnerabilities, we can understand how a defensive mechanism has been adopted. On the other hand, Internet wide security scanning means that we have to probe the entire address space and with existing tools, this is a very difficult and slow task. Some of the security studies does require scanning of the entire

address space. For example, if we can obtain the cryptographic keys in the entire address space, then we can understand the ecosystem of public key infrastructure and its vulnerabilities.



Electronic Frontier Foundation on public key certificates. Again, the study requires a lot of effort and a comprehensive scan of the whole Internet also requires months of efforts. All these projects require heroic effort just to perform initial data collection. As you can see, these scans require massive parallelization and typically they also require an extended period of time. Given the benefits of Internet-wide network security studies, we want to see a few more of them.



But in order to do that, Internet surveys should not require heroic efforts. For example, it would be nice if we can scan the whole HTTPS ecosystem every day. Given that existing tools are so slow, researchers at the University of Michigan developed ZMap, a new Internet scanner, from scratch. ZMap is a fully functional open-source network scanner that can scan a single port on the IPv4 address space. On a gigabit network, it can cover 98 percent of the IPv4 address space from a single machine in under 45 minutes.

With Zmap, an Internet-wide TCP SYN scan on port 443 is as easy as:

```
$ zmap -p 443 -o results.txt
34,132,693 listening hosts
(took 44m12s)
```

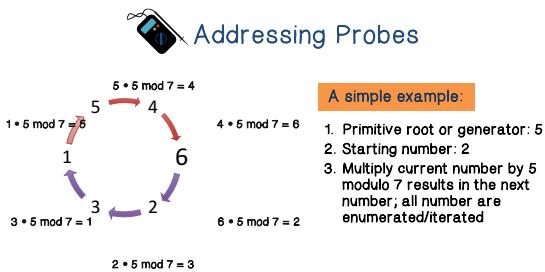
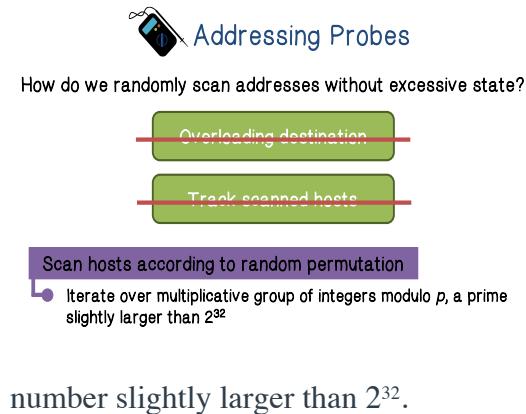
97% of gigabit Ethernet linespeed

With ZMap, performing a scan on the Internet no longer requires hundreds of hours of coding or months to execute. Now, performing an Internet scan is a single command that could be finished within an hour.

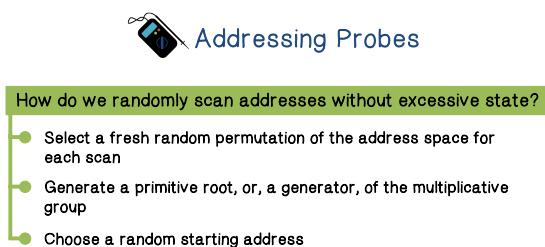
ZMap Architecture	
Existing Network Scanners	ZMap
Reduce state by scanning in batches <ul style="list-style-type: none"> • Time lost due to blocking • Results lost due to timeouts 	Eliminate local per-connection state <ul style="list-style-type: none"> • Full asynchronous components • No blocking except for network
Track individual hosts and retransmit <ul style="list-style-type: none"> • Most hosts will not respond 	Shotgun Scanning Approach <ul style="list-style-type: none"> • Always send n probes per host
Avoid flooding through timing <ul style="list-style-type: none"> • Time lost waiting 	Scan widely dispersed targets <ul style="list-style-type: none"> • Send as fast as network allows
Utilize existing OS network stack <ul style="list-style-type: none"> • Not optimized for immense number of connections 	Probe-optimized Network Stack <ul style="list-style-type: none"> • Bypass inefficiencies by generating Ethernet frames

While there are many excellent multi-purpose network scanners, these scanners were never intended or optimized for scanning the entire Internet. ZMap, on the other hand, was designed from the ground up for the specific purpose of

completely scanning the whole Internet. Let us compare existing network scanners with ZMap. Existing network scanners need quite an immense amount of resources to keep track of state of the entire Internet. ZMap, on the other hand, eliminates local, per-connection state. Existing scanners also spend a lot of resources keeping track of which host is being scanned and which host have responded. On the other hand, ZMap uses a shotgun scanning approach and keeps minimum state information. Previous scanners have been attempting to be more polite by slowing down the scans. ZMap, on the other hand, distributes the scan across the Internet. This allows ZMap to scan at more aggressive rate without impacting the destination networks. Finally, ZMap also bypasses the inefficient network stack in the operating system and generates outgoing packets directly.

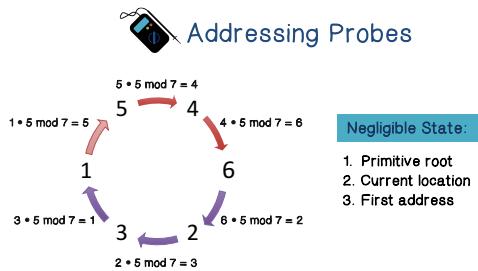


If we simply poll network addresses in numerical order, we will risk overloading the destination networks. However, we also do not want to have to track what host we have scanned or need to scan. In order to scan according to a random permutation, we select addresses according to a cyclic group. Specifically, we iterate over a multiplicative group of integers modulo a prime



Here is an example of how we randomly iterate through all the numbers in a group. The generator is 5, and we start with number 2. The process is that we multiply the current number by 5 mod 7, and the result is the next number, and all numbers will be enumerated. For example, we start with 2. 2 multiply by 5 mod 7 is 3. So, we move to 3. 3 times 5 mod 7 is 1, and so on. As you can see, we can iterate through all the numbers in a group.

Following that simple example, here is how we can iterate through the addresses in address space. Each scan is essentially a random permutation of the address space. Again, we treat the address space as a multiplicative group. So, we first decide on the generator and then we choose a random starting address.



addresses in the address space. As you can see, this is very little information to keep track of for each scan.

TCP IP Quiz

Step 1: Which protocol is used to break data into packets?

TCP

Step 2: Which protocol is used to move packets from router to router?

IP

Step 3: Which protocol reassembles the data packets?

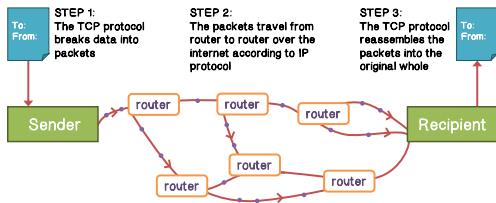
TCP

With this approach, we can guarantee that all addresses in the address space can be iterated. More importantly, there is very little state information we need to keep track of. Essentially, for each scan, we need to keep track of the primitive root or the generator and the current address because when the first address is visited again, we know that we have iterated all the

Now, let us do a quiz. Which protocol is used to break data into packets? It is TCP. Which protocol is used to move packets from router to router? It is IP, the Internet Protocol. Which protocol reassembles the data packets? Again, it is TCP.

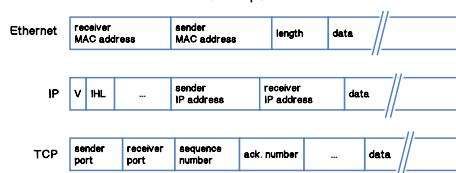
TCP IP Quiz

How TCP/IP Works



Validating Responses

How do we validate responses without local per-target state?
Encode secrets into mutable fields of probe packets that will have recognizable effect on responses

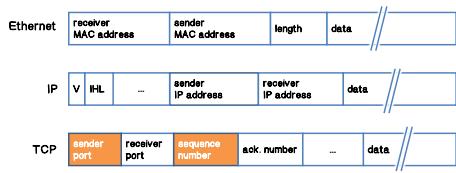


And here is an illustration of the concepts in the quiz. On the sender's side, the TCP protocol breaks data into packets. When a packet comes out of the sender, the routers are responsible for sending a packet to the recipient. At the recipient the TCP protocol again reassembles the packets back into the original data.

We have discussed how we send probe packets without storing what hosts we have scanned or need to scan, but how do we validate the responses? The idea here is that we can encode secrets into the fields of the probe packets. This is similar to the SYN cookies. Here are the formats of the Ethernet, IP and TCP packets, and some of the fields here can be used to encode the secrets.

✓ Validating Responses

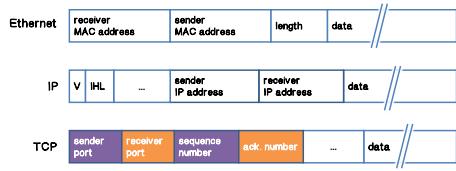
How do we validate responses without local per-target state?
Encode secrets into mutable fields of probe packets that will have recognizable effect on responses



using a TCP SYN packet for scanning, we encode the secret into the source port and sequence number. We know that the destination host will have to include the source address and port and acknowledge the sequence number to send a response.

✓ Validating Responses

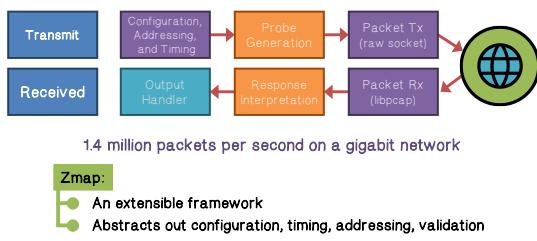
How do we validate responses without local per-target state?
Encode secrets into mutable fields of probe packets that will have recognizable effect on responses



In order to validate the packets that we receive, we encode a scan invocation and host-specific secret into mutable fields that will have a recognizable effect on the probe responses. Specifically, for each host to be scanned, ZMap computes a hash of the destination address keyed by a scan-specific secret. This MAC value, or Message Authentication Code, is then spread across any appropriate and available fields. In the case of

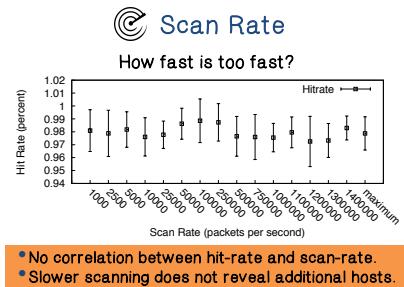
For example, a sequence number will be encoded into the acknowledgement number field. More precisely, in the response packet the receiver port was the sender port in the scanning packet and the acknowledgement number is the sequence number of the scanning packet.

Wi-Fi Packet Transmission and Receipt

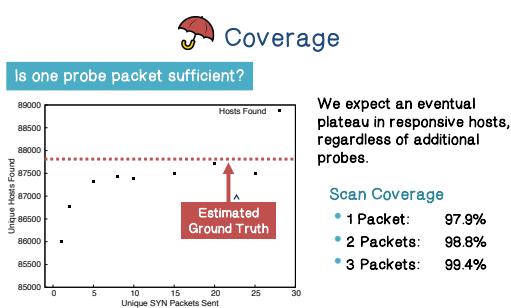


Here are the high-level workflows of ZMap. To send scanning packets, ZMap will be configured, for example, with a random permutation of the address space. And then the probing packets will be sent to these addresses according to the random permutation. Since the majority of fields in a probe packet never change, ZMap performs all network operations at the Ethernet layer using a raw socket in order to cache values. This eliminates time lost

to kernel operations such as route lookups. This allows ZMap to send probing packets at a very high speed and validate the responses. For example, checking the MAC value and the probing results can then be analyzed. This configuration allows ZMap to send probes at approximately 1.4 million packets per second on a gigabit network. ZMap is an extensible framework to allow any type of single packet scan such as a TCP SYN scan, ICMP echo requests or application specific UDP scan. The ZMap framework abstracts out details such as configuration, timing, addressing and validation.



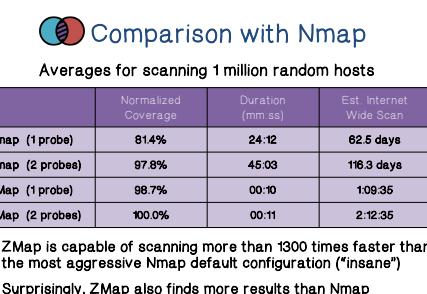
card. As we can see there is no correlation between the scanning rate and the number of hosts that respond. One way to interpret this result is that even slower scanning rates does not produce additional response.



number of responsive hosts regardless of additional probes. We can estimate that with a single packet the coverage is about 97.9 percent and with three packets the coverage is 99.4 percent.

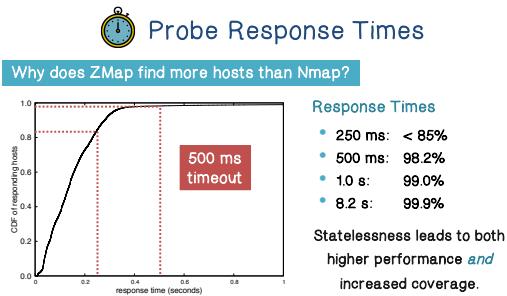
The first major question is whether ZMap or our network can handle scanning at gigabit speed. In order to answer this question researchers at the University of Michigan performed 10 trials scanning one percent of the IPv4 address space at different rates ranging from 1,000 packets per second to the maximum 1.4 million packets per second. This maximum is determined by the NIC configuration. NIC stands for network interface

Another interesting question is, “Is one probe packet sufficient or is sending multiple packets beneficial?” This is difficult to answer directly because there is no ground truth for the number of live hosts on the Internet. In order to estimate ZMap coverage, multiples packets were sent to one percent samples of the IPv4 address space. If you look at the number of responses, we clearly see a plateau after a one-SYN packet scan. In fact, we should expect to see an eventual plateau in the



these results, ZMap is capable of scanning the IPv4 address space more than 1,300 times faster than Nmap. ZMap also has higher coverage than Nmap even when Nmap sends multiple probes. To be fair, Nmap is an excellent multi-purpose network scanner and it is optimized for completely different use cases than ZMap.

Many previous projects have used the popular Nmap network scanner to perform scans. Therefore, we should compare ZMap with Nmap. The researchers performed several experiments focusing on the total time spent for scanned and scan coverage. In these experiments, one million addresses were scanned. In these experiments, Nmap used the most aggressive scan template called *insane* and with a minimal scan rate of 10,000 packets per second. As we can see from



both increased performance and increased coverage.

So why does ZMap have higher coverage than Nmap? We analyzed the response time of the scan responses. ZMap does not timeout hosts, but Nmap does. In fact, if Nmap sends one packet, it times out after 250 milliseconds. If it uses two packets, it times out after 500 milliseconds. As you can see, some responses arrive after Nmap has timed out. Ultimately, since ZMap uses stateless scanning and does not keep state, it has

Entropy Quiz

Fill in the blanks with the correct answers:

With regards to computing, what is entropy?

Randomness for use in cryptography or other applications that require random data.

What are the two sources of entropy?

Hardware sources and randomness generators

A lack of entropy will have a **negative** impact on performance and security.

Now let us do a quiz on entropy. Fill in the blanks with the correct answers.

With regards to computing. What is entropy? Entropy is randomness for use in cryptography or other applications that require random data. What are the two sources of entropy? Hardware sources or randomness generators. A lack of entropy will have a negative impact on performance and security.

Cryptographic Keys

Uncovering weak cryptographic keys and poor entropy collection

	HTTPS	SSH
Live Hosts	12,8 million	10,2 million
Distinct RSA Public Keys	5,6 million	3,8 million
Distinct DSA Public Keys	6.241	2,8 million

Now, let's discuss a few interesting Internet wide security studies using ZMap. In one study, the researchers scanned the HTTPS servers and the SSH servers and collected their public keys. As you can see, a large number of machines share their public keys. Although there are many legitimate reasons why machines would share their keys, we need to understand whether the keys happen to be the same without the intention of sharing because that is bad for security.

Cryptographic Keys

Why are a large number of hosts sharing cryptographic keys?

5.6% of TLS hosts and 9.6% of SSH hosts share keys in a vulnerable manner

- Default certificates and keys
- Apparent entropy problems

In particular, researchers find that some of these machines share their key in a vulnerable manner. For example, they simply use the default keys without creating their own keys and there is an apparent entropy problem that caused them to have the same keys.

🔑 Factoring RSA Public Keys

What else could go wrong if devices aren't collecting entropy?

RSA Public Key: $n = p \cdot q$, p and q are two large random primes

Most efficient known method of compromising an RSA key is to factor n back to p and q

While n is difficult to factor, for

$N_1 = p \cdot q_1$ and $N_2 = p \cdot q_2$

we can trivially compute

$p = \text{GCD}(N_1, N_2)$

(GCD) of N_1 and N_2 . The reason that two different public keys may have the same prime number can be that there is not sufficient entropy and therefore two machines may happen to generate the same large prime number p . There is also a very efficient algorithm due to D. J. Bernstein to compute the GCD of every pair of N_1 and N_2 .

Let us look at the effect of not having sufficient entropy. As a quick review, in RSA, the public key n is the product of p and q , where p and q are large random prime numbers. The security of RSA is based on the fact that factoring n back to p and q is very, very inefficient. That is, it is very hard for the attacker. On the other hand, if two private keys N_1 and N_2 share the same prime number, say p , then it is trivial to compute p because p is the greatest common denominator

🔑 Factoring RSA Public Keys

Why are a large number of hosts sharing cryptographic keys?

- We find 2,134 distinct primes and compute the RSA private keys for 64,081 (0.50%) of TLS hosts
- Using a similar approach for DSA, we are able to compute the private keys for 105,728 (10%) of SSH hosts
- Compromised keys are generated by headless or embedded network devices
- Identified devices from > 40 manufacturers



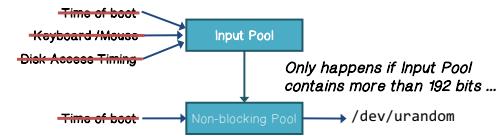
With this efficient approach, the researchers were able to crack many keys on TLS and SSH machines. The majority of these machines are network devices.

🔑 Embedded Systems

Linux /dev/urandom
Nearly everything uses /dev/urandom

Problem 1 Embedded devices may lack all these sources

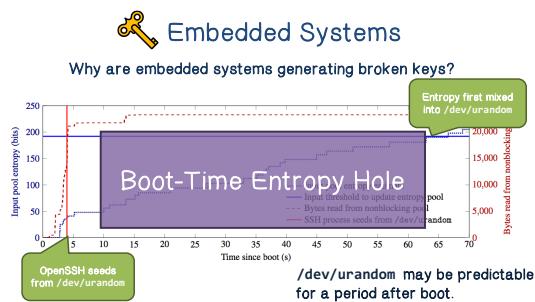
Problem 2 /dev/urandom can take a long time to "warm up"



Now, let us look at why these embedded systems or network devices generate broken keys. These embedded systems all run Linux and they all use urandom to generate cryptographic keys. Linux maintains several entropy pools. As entropy is gathered, it is stored in the input pool and is eventually mixed into the non-blocking pool which feeds urandom. However, many of these sources are not present on an embedded system. For example, the system may not have keyboard

and many of them do not have spindle-based disks. In fact, most of these embedded systems do

not even support real-time clocks. Therefore, with all of these random sources removed, we end up with a deterministic source of randomness. Furthermore, entropy is a mix from the input pool to the non-blocking pool until there are more than 192 bits of entropy that have been collected. Therefore, even if a small amount of entropy has been collected on a device, none of it is available until the pool reaches this threshold.



predictable urandom.

This graph shows the build-up of entropy on the first boot of a typical Ubuntu desktop server. No randomness is added from the input pool to urandom until 192 bits of entropy is available. From the graph, we can tell that this did not happen until fairly late in the boot process. This is the so-called Boot-Time Entropy Hole problem, which means that urandom may be predictable for a period after boot. Unfortunately, cryptographic keys may be based on this



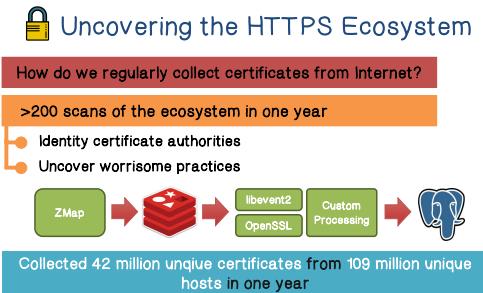
them.



to sign Google with authority. And this can go on involving multiple intermediates and we end up having a certificate chain. That is, if you look at Google.com, it is signed by a chain of certificates and a top certificate is just a self-signed certificate. This entire chain is presented by Web browser to the client and the client can verify the signatures by traversing the certificate chain up to the browser trusted certificate, which is self-signed.

Another interesting study is on the certificates issued by certificate authorities. This is an important topic because HTTPS underlines all secure Web communications and HTTPS is depending on the security of certificate authorities. The problem is that there are many certificate authorities and every one of them can sign for any website. In fact, we do not even know all the certificate authorities, until we see

Let us have a quick review of how certificates are used in Web browsing. Browsers such as Firefox decide whom they trust. These certificates are stored as part of the browser or the operating system. Browsers typically support a couple hundred root certificates, such as shown in this figure. These root certificate authorities then agree to sign other certificates. For example, Equifax is a root certificate authority and agrees



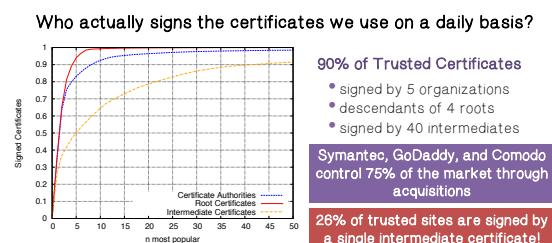
The researchers at the University of Michigan performed regular scans to track the use of certificates. They observed 3,700 browser trusted certificates in one year. They also discovered two cases of misused certificates. In one case, a signing certificate was accidentally issued to a Turkish transit provider. In another case, 1,300 certificates were issued by the Korean government.

⌚ Identifying Certificate Authorities



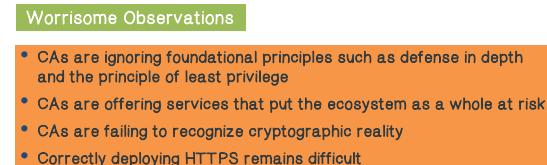
The researchers found 1,800 signed certificates belonging to 683 organizations and these include various kinds of institutions. The most worrisome finding is that all major browser roots are selling intermediates to third parties without any constraints.

⌚ Identifying Certificate Authorities



By analyzing the certificate chains, we see that 90 percent of the certificates are signed by five organizations. They are sent from four roots and signed by 40 intermediates and there are only a few big players. Another interesting fact is that 26 percent of the trusted sites are signed by only a single intermediate certificate.

⚠ CA Risks



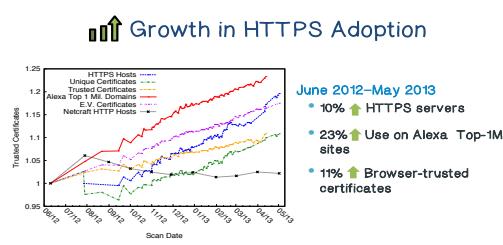
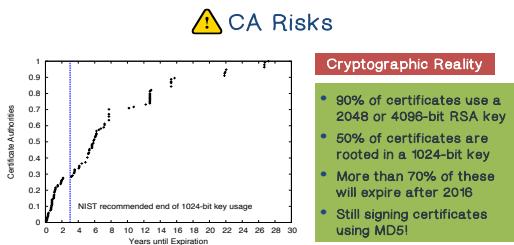
At a high level, there are several worrisome observations. First, the certificate authorities are ignoring defense in depth and least privilege. They are offering services that put the whole ecosystem at risk and they are using some weak keys and deploying HTTPS remains difficult. Let us go on with these observations next.

⚠ CA Risks

Ignoring Foundational Principles

- We classically teach concepts such as defense in depth and the principle of least privilege
- We have methods of constraining what CAs can sign for, yet all but 7 of the 1,800 CA certs we found can sign for anything
- Lack of constraints allowed a rogue CA certificate in 2012, but in another case prevented 1,400 invalid certificates
- Almost 5% of certificates include local domains *e.g.* localhost, mail, exchange

names are not fully qualified and intended resource is ambiguous and there is no identifiable owner. As such, these local domain names frequently appear on more than one certificate. In one example, there are about 1,218 browser trusted certificates for the domain mail owned by organizations ranging from the U.S. Department of Defense to small companies.



ZMap Open Source

Releasing ZMap as a fully documented open source project

Downloaded it from <https://zmap.io>

Scanning the Internet *really* is as simple as:

```
$ zmap -p 443 -o results.csv
```

Be sure you have adequate bandwidth and be a good Internet neighbor!

Scans.IO Data Repository

How do we share all this scan data?

University of Michigan is hosting a repository of data gathered from Internet-wide scans:
<https://scans.io>

For defense in depth, there are several technical practices already in place for limiting the scope of a signing certificate including setting name or path length constraints and distributing leave certificates among a large number of intermediate certificates. There are clear cases for using these restrictions but the vast majority of the time, the CAs do not utilize these options. As another example, local domain

In terms of the keys used, the researchers also found that 90 percent of the certificates use a 2048-bit or 4096-bit RSA keys. But 50 percent of the certificates were rooted in the 1024 bit key. More importantly, more than 70 percent of these would expire after 2016 and many still signed using MD5.

The researchers also found that within a year, the number of HTTPS servers only increased by 10% and there was a 23% increase in the number of Alexa top 1 million sites using the HTTPS. And there was 11% increase in the raw number of browser trusted certificates. Here is an example of how we use Internet wide scan to check the adoption of a technology.

The researchers at the University of Michigan had released ZMap as an open source project. You can download ZMap and use it for your own Internet wide security study. Of course, please remember to be a good citizen on the Internet.

There is also a repository of data from previous scans.

ZMap: Fast Internet-wide Scanning and Its Security Applications

Zmap is a modular single packet network scanner designed for interne-wide network surveys. On a computer with a gigabit connection, Zmap can scan the entire public IPv4 address space in under 45 minutes. Although scanning is usually associated with botnets and worms, it can also be used as a valuable methodology in security research. It can be used to reveal all kinds of vulnerabilities, monitor deployments of mitigations and throw light on previously opaque distributed ecosystems.

The paper starts by giving an introduction to Zmap and all of its advantages, it compares Zmap to Nmap, which is a widely used network mapping tool. Zmap performs much better than Nmap for the following reasons (page 1 and page 2):

1. Optimized probing
2. No per-connection states
3. No retransmission

Section 2 describes the working of Zmap. It starts by explaining how the scan addresses are decided using randomized permutations of the address space and how it excludes certain address spaces. The packet transmission and receipt are designed in such a way that it maximizes the number of probes it can send based on the source's CPU and NIC. The probe modules are responsible for filling in the body of the probe and validate the incoming packets.

This section is followed by a set of experiments performed on Zmap and their results. It compares if the scan rate of Zmap has any effects on the hit rate and checks if coverage is an issue at such speeds (Section 3.1 and 3.2). Then there is an extensive comparison of Zmap with Nmap and other previous methods that have been used for Internet scanning. The comparisons include time taken, the coverage and timeouts (Section 3.4 and 3.5).

The next section explores the research possibilities with the ability to scan the entire IPv4 address space. This can be used for increased visibility into distributed systems, tracking how successful the adoption of new protocols is, enumerating vulnerable hosts and discovering unadvertised services that were previously available only with explicit knowledge of host name or address (Sections 4.1 to 4.5). This is followed by suggestions for other researchers who are performing fast Internet wide scans, these are guidelines that ensure that these tools are not used with the wrong intentions.